

Lab5 – Knowledge Distillation

Advisor : Tsai, Chia-Chi

TA : 柯秉鈞

Contact : course.aislab@gmail.com

Outline



1. Knowledge Distillation
 - Introduction
 - Types of Knowledge
 - Modes of Distillation
2. Lab5 task

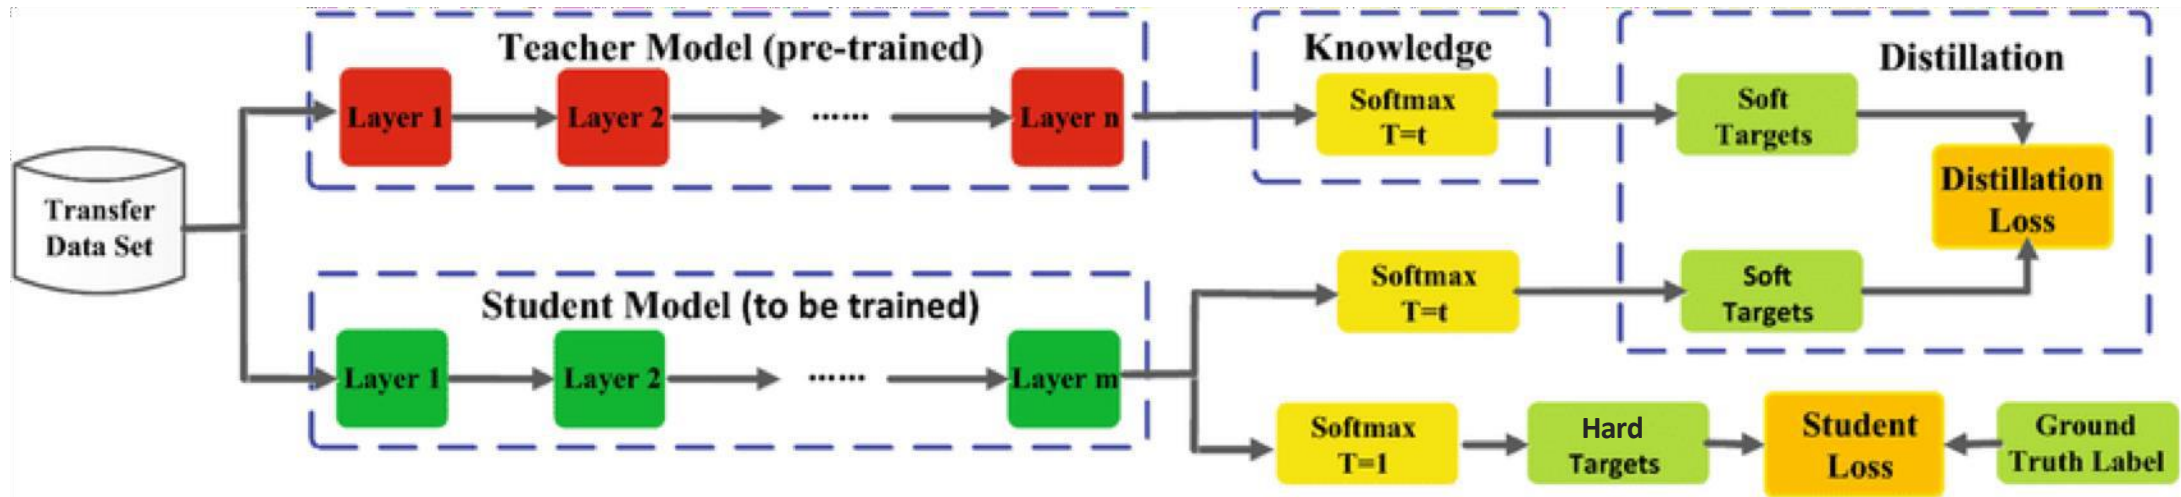
Outline



1. Knowledge Distillation
 - Introduction
 - Types of Knowledge
 - Modes of Distillation
2. Lab5 task

Knowledge Distillation

- A small model (student) is trained to mimic a large pre-trained model (teacher)



$$q_i = \frac{\exp(z_i/T)}{\sum_j \exp(z_j/T)}$$


Temperature



Knowledge Distillation

- Temperature for softmax

$$y'_i = \frac{\exp(y_i)}{\sum_j \exp(y_j)} \xrightarrow{T=100} y'_i = \frac{\exp(y_i/T)}{\sum_j \exp(y_j/T)}$$



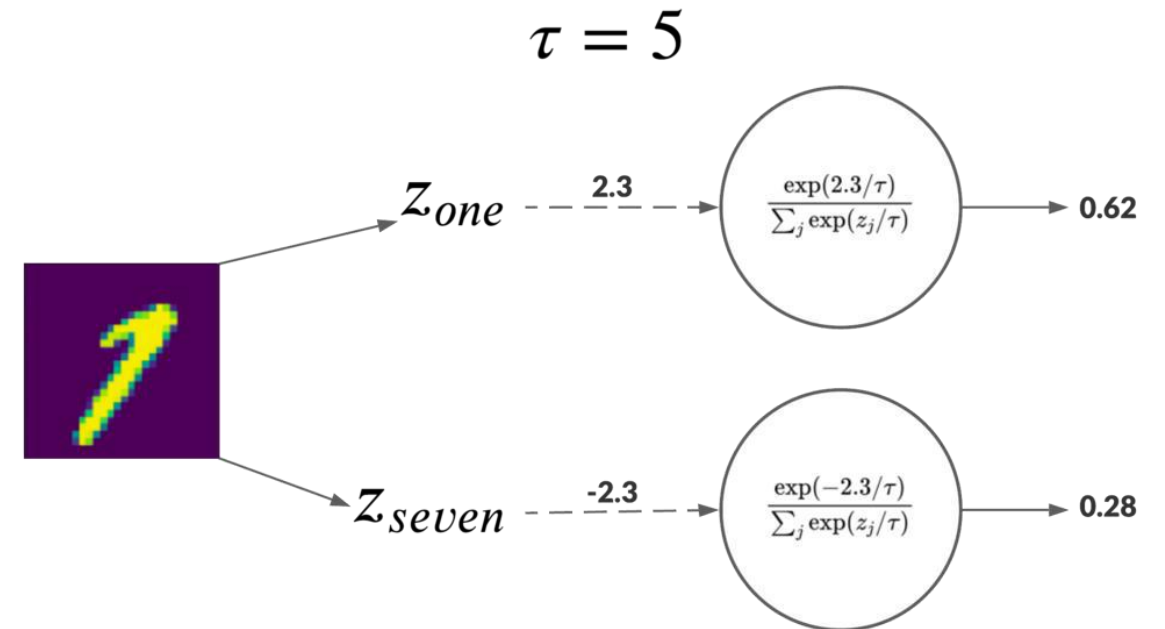
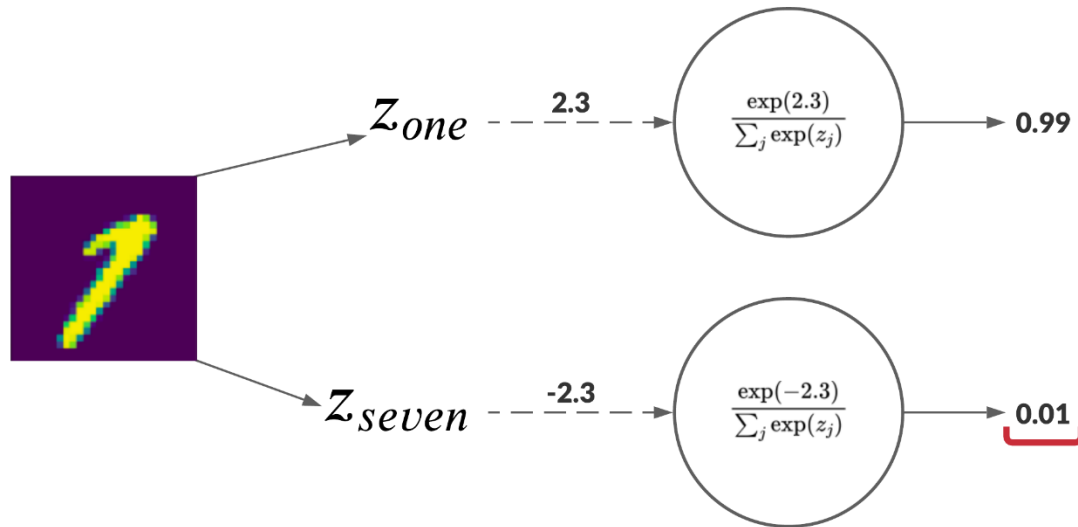
$y_1 = 100$	$y'_1 = 1$
$y_2 = 10$	$y'_2 \approx 0$
$y_3 = 1$	$y'_3 \approx 0$



$y_1/T = 1$	$y'_1 = 0.56$
$y_2/T = 0.1$	$y'_2 = 0.23$
$y_3/T = 0.01$	$y'_3 = 0.21$

Soft Target

- Model prediction softened by T

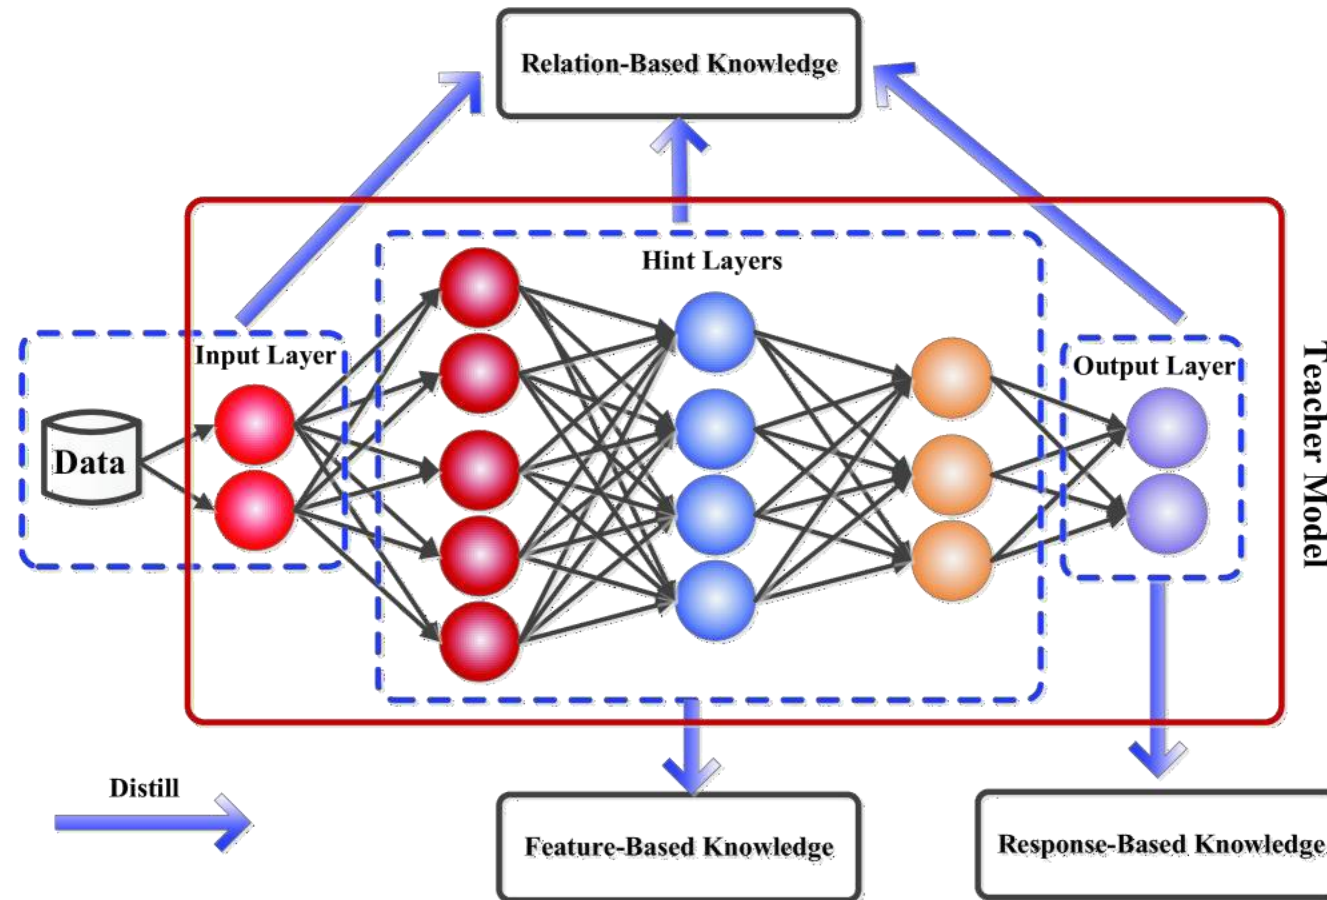


Outline

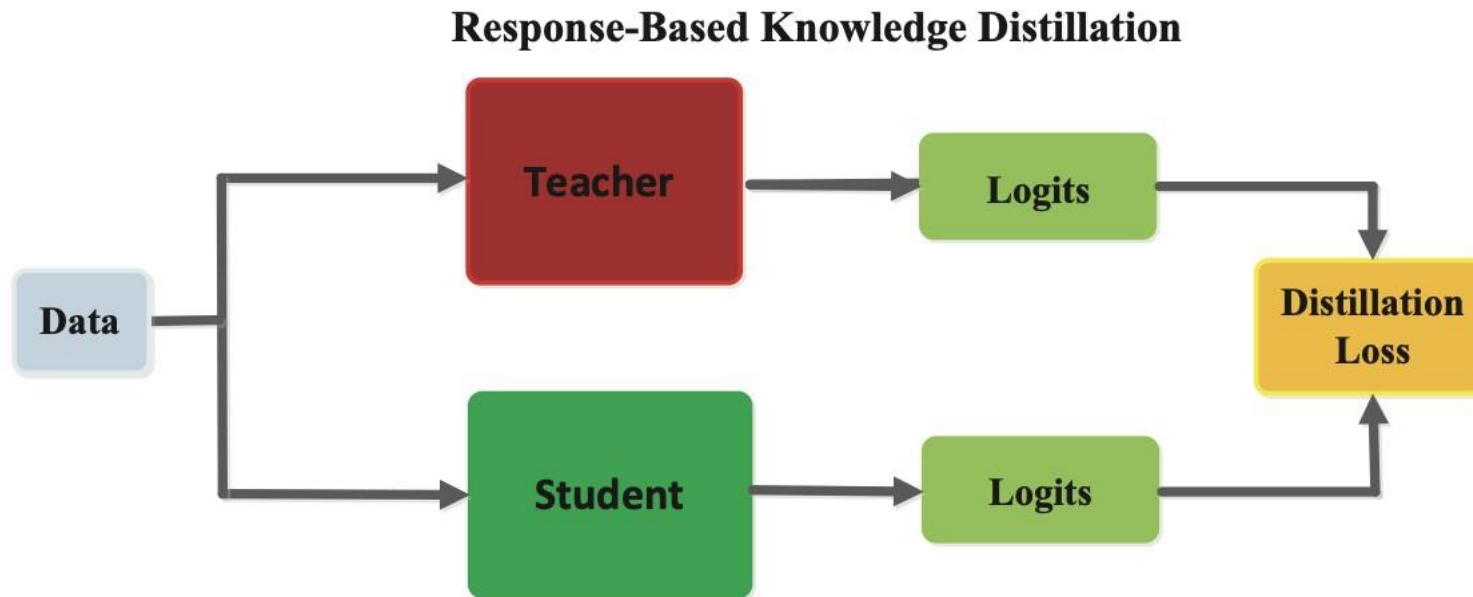


1. Knowledge Distillation
 - Introduction
 - Types of Knowledge
 - Modes of Distillation
2. Lab5 task

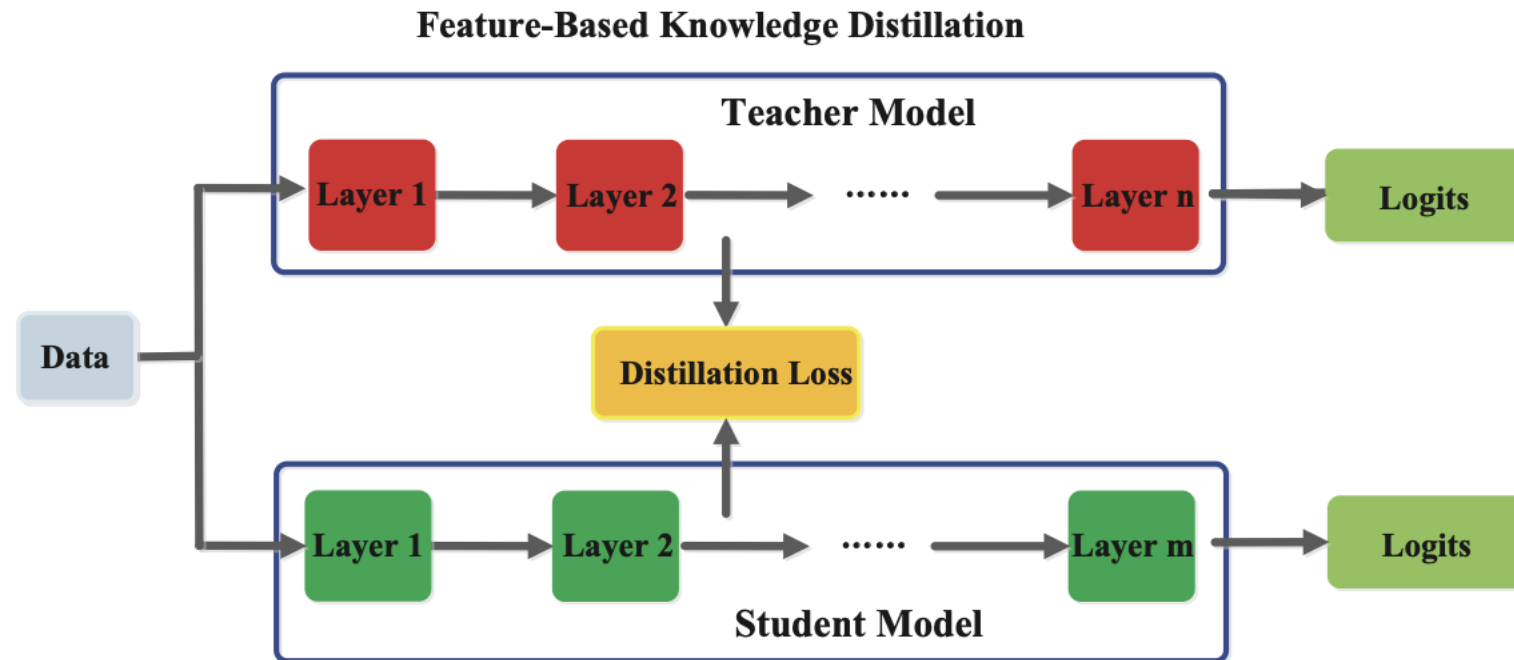
Types of Knowledge



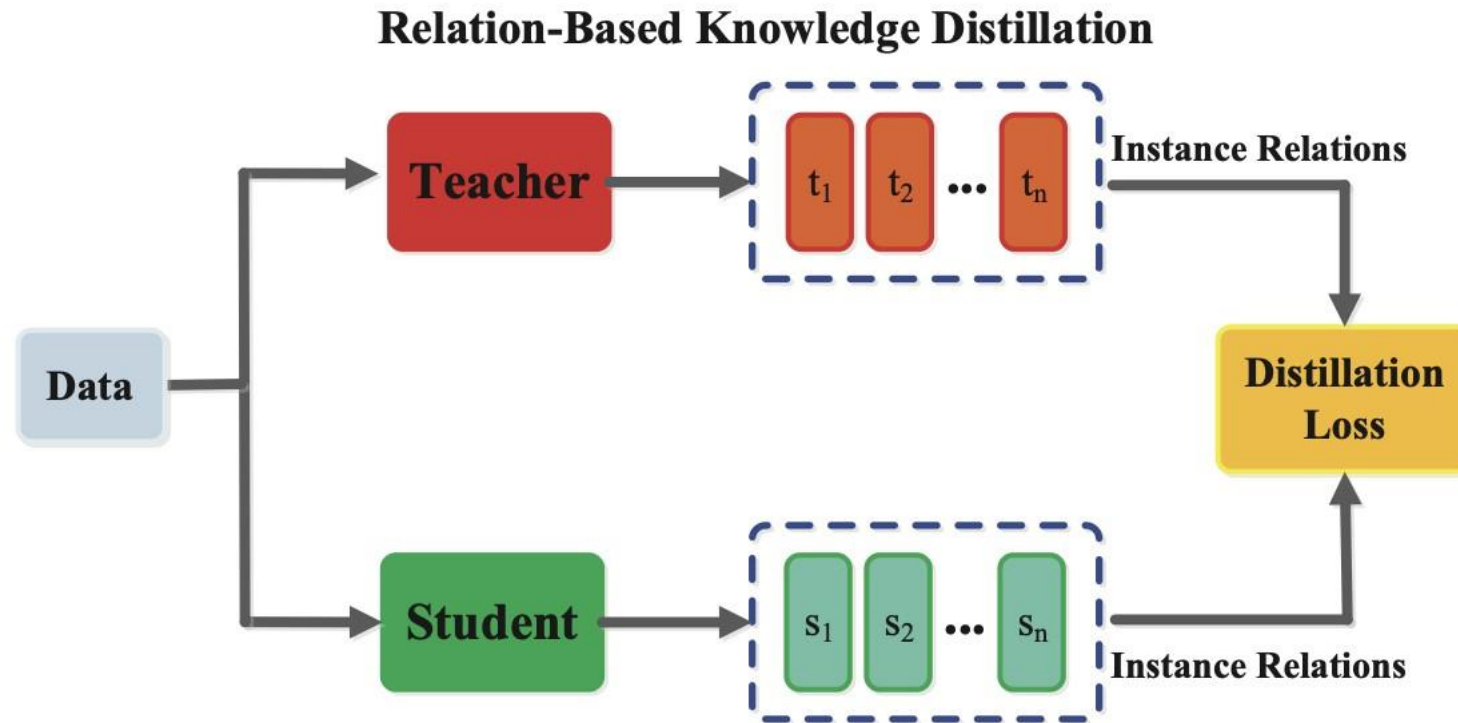
Response-Based Knowledge



Feature-Based Knowledge



Relation-Based Knowledge

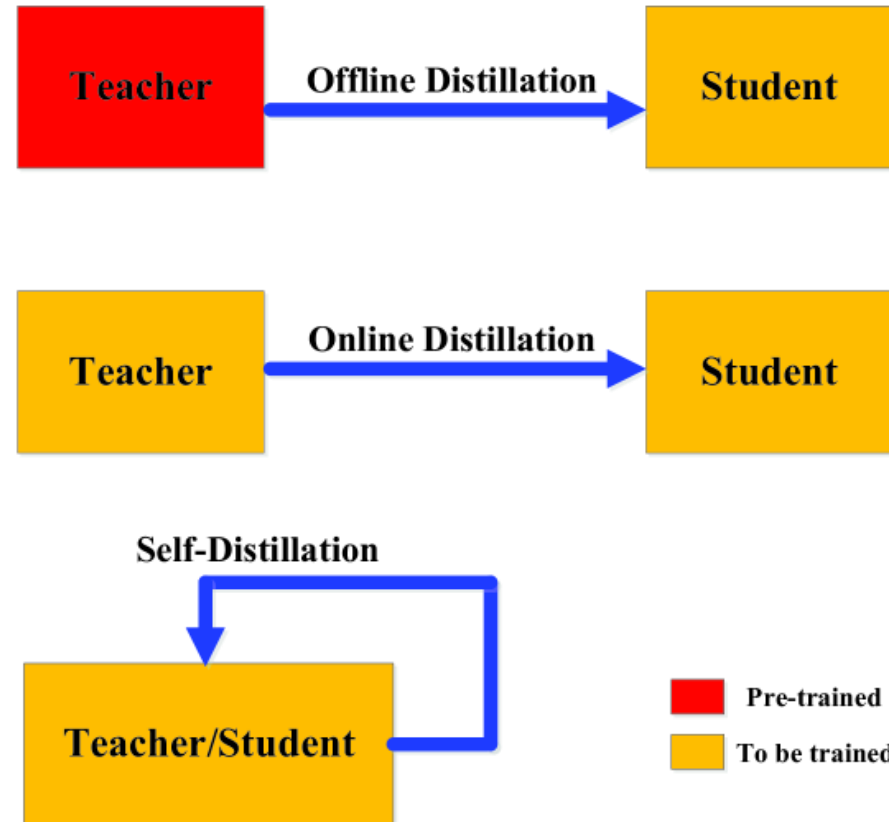


Outline



1. Knowledge Distillation
 - Introduction
 - Types of Knowledge
 - Modes of Distillation
2. Lab5 task

Modes of Distillation



Outline



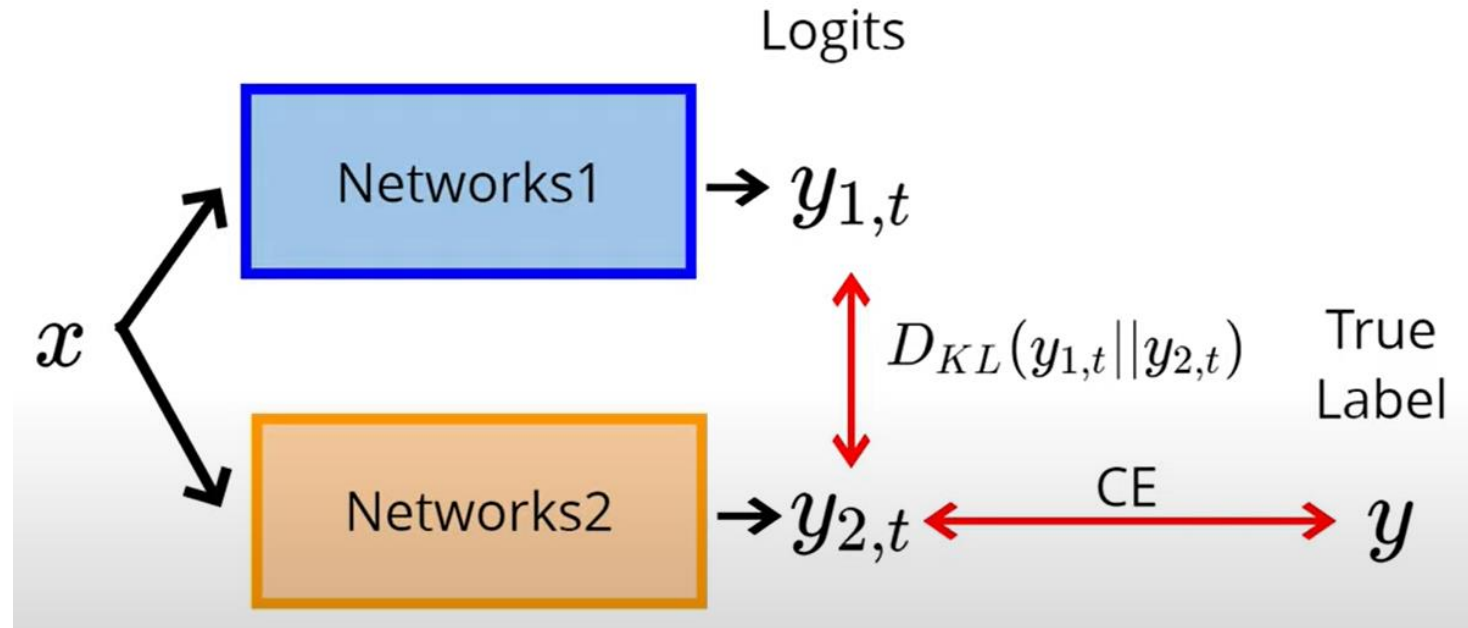
1. Knowledge Distillation
 - Introduction
 - Types of Knowledge
 - Modes of Distillation
2. Lab5 task

Lab5 task



- Teacher: ResNet34
- Student: ResNet18
- Dataset: Cifar10
 - Response-based distillation (**35%**, **accuracy ≥ 0.8 got 30%**, the more improvement the higher score **compare with student from scratch**)
 - Feature-based distillation (**35%**, same as above)
 - Report (**30%**)
- [Colab link](#)
- **StudentID_lab5.ipynb**

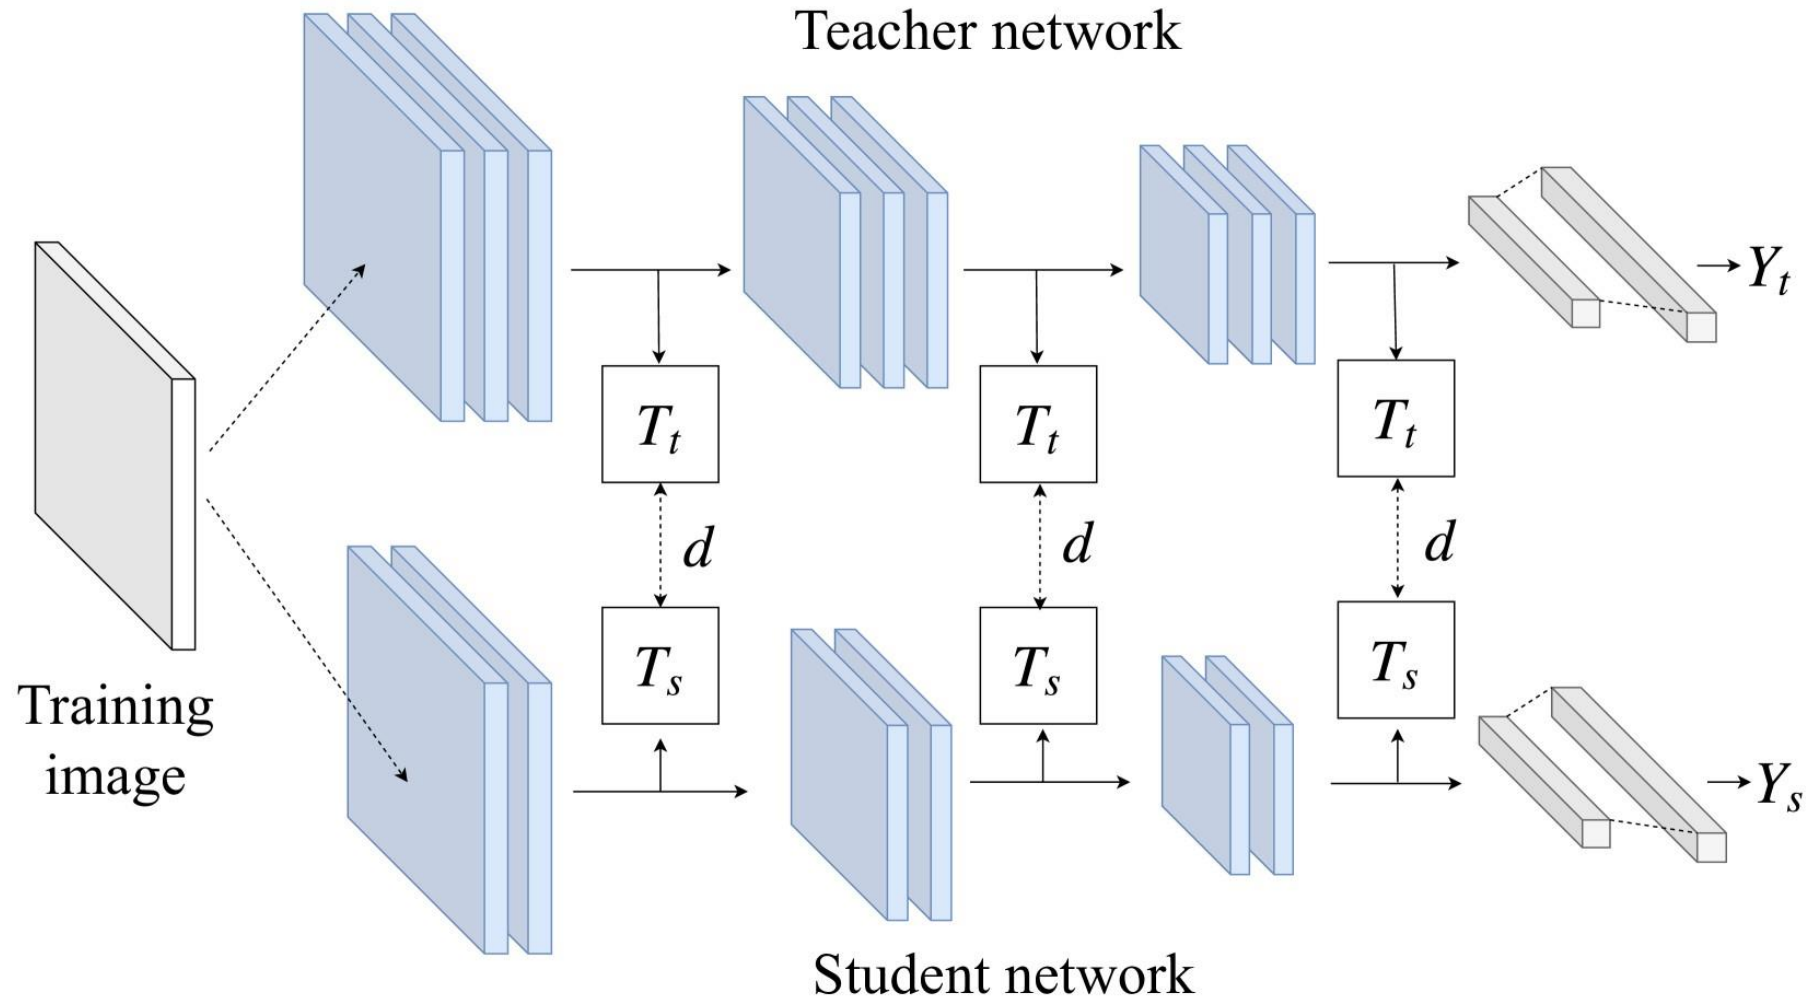
Response-Based Knowledge



$$q_i = \frac{\exp(z_i/T)}{\sum_j \exp(z_j/T)}$$

$$L = \alpha T^2 L^{(soft)} + (1 - \alpha) L^{(hard)}$$

Feature-Based Knowledge



Reference



- [\[2006.05525\] Knowledge Distillation: A Survey \(arxiv.org\)](#)
- [\[1904.01866\] A Comprehensive Overhaul of Feature Distillation \(arxiv.org\)](#)
- [\[1503.02531\] Distilling the Knowledge in a Neural Network\(arxiv.org\)](#)

Thanks for listening